

Pipes

COMS10012 Software Tools

standard IO



Unix Philosophy

It is easier to maintain 10 small programs than one large program. Therefore,

1. Each program should do one thing well.
2. Programs should be able to cooperate to perform larger tasks.
3. The universal interface between programs should be a text stream.

source

```
#include <stdio.h>
// gives stdin etc.
// fread, fwrite, FILE* - C abstraction

#include <unistd.h>
// pulls in /usr/include/sys/unistd.h
// read, write - POSIX abstraction

#define STDIN_FILENO    0
#define STDOUT_FILENO   1
#define STDERR_FILENO   2
```

standard input/output

Internally, programs `read(fd, buffer, size)` and `write(fd, buffer, size)`.

Each program starts with three file descriptors open:

0 = standard input

1 = standard output

2 = standard error



standard input/output

Running a program in the terminal:

■ program ■



shell



pipes




pipe

```
$ ls -1 | head
```

```
head [-n NUM]
```

```
tail [-n NUM]
```



pipe

```
$ ls -1 | grep software | sort
```

grep: "global regular expression parser"

sort: read all lines into buffer, sort, output

uniq: remove duplicates immediately following

best used as: `command | sort | uniq`

grep

```
$ grep PATTERN FILENAMES
```

```
$ grep -nHi PATTERN FILENAMES
```

```
$ grep [OPTIONS] PATTERN
```



sort

\$ sort

aaa

ccc

bbb

^D

aaa

bbb

ccc

\$

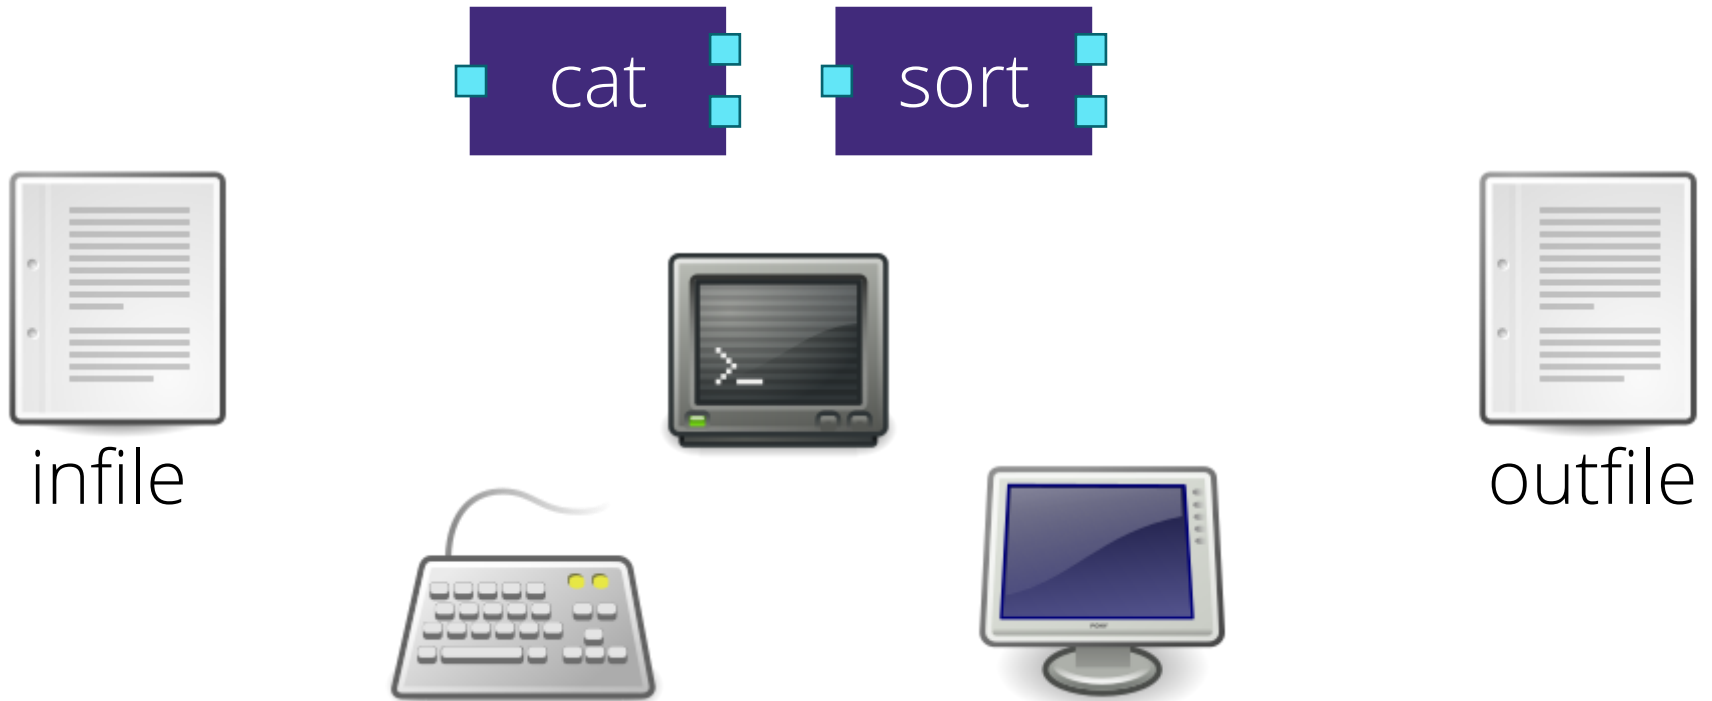


redirects



redirect

```
$ cat infile | sort > outfile
```



redirect

```
$ sort < infile > outfile
```



infile



outfile



redirect

\$ COMMAND > FILE

overwrites FILE

\$ COMMAND >> FILE

appends to FILE

error redirect

```
$ COMMAND > FILE 2> FILE2
```

```
$ COMMAND > FILE 2>&1
```

not:

```
$ COMMAND 2>&1 > FILE
```

ignore output:

```
$ COMMAND > /dev/null
```



files vs streams

A program that uses a standard stream can be told to use a file instead by

- `PROGRAM < FILE` (standard input)
- `PROGRAM > FILE` (standard output)
- `PROGRAM 2> FILE` (standard error)

files vs streams

A program that expects a filename can be told to use standard input/output instead by:

- using the filename `-` (single dash), if the program supports it
- using the filename `/dev/stdin` etc., if your OS supports it



Filenames with dashes

Filenames starting with dashes are generally considered bad.

If you really want to address one (e.g. you created one by mistake), use e.g.

```
$ cat ./-
```

```
$ rm ./-f
```

advanced



tee

```
$ ls | tee FILE
```

tee: takes a filename
as argument and writes
a *copy* of input to it,
as well as to stdout



need a file, want a pipe

If PROGRAM wants a file to read from, how can I pipe something in?

```
$ PROGRAM <(SOMETHING)
```

```
$ cat <(echo "Hi")
```

Hi

```
$ echo <(echo "Hi")
```

```
/dev/fd/63
```



subshell

```
$ cat <(echo "Hi")
```

echo



/dev/fd/...

cat



subshell to argument

```
$ COMMAND $(SOMETHING)
```

```
$ echo $(echo Hi | sed -e s/Hi/Hello/)  
Hello
```

old-fashioned way, with backticks:

```
$ COMMAND `SOMETHING`
```



